

Phase 3: Reinforcement Learning & AI Integration

Advanced ML-Powered Trading Enhancement

Enhance your SMC Trading Bot with Reinforcement Learning and AI-assisted pipelines using free-tier services.

Overview

Phase 3 adds:

- **Deep Q-Network (DQN)** Reinforcement Learning agent
- **AI-assisted signal generation** pipeline
- **Continuous learning** from trade outcomes
- **Google Colab** integration for free GPU training
- **Enhanced dashboard** with ML metrics

Part 1: RL Agent Implementation

Step 1: Create RL Agent Module

Create backend/rl_agent.py:

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.models import Sequential, load_model, save_model
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.optimizers import Adam
from collections import deque
import random
import os

class DQNAgent:
    """
    Deep Q-Network Reinforcement Learning Agent
    Learns optimal trading actions from experience
    """

    def __init__(self, state_size=50, action_size=3, learning_rate=0.001):
        self.state_size = state_size
        self.action_size = action_size # 0: HOLD, 1: BUY, 2: SELL

        # Hyperparameters
        self.memory = deque(maxlen=2000)
        self.gamma = 0.95 # Discount factor
```

```

        self.epsilon = 1.0           # Exploration rate
        self.epsilon_min = 0.01
        self.epsilon_decay = 0.995
        self.learning_rate = learning_rate
        self.batch_size = 32
        self.train_start = 100

        # Build neural networks
        self.model = self._build_model()
        self.target_model = self._build_model()
        self.update_target_model()

        self.training_step = 0

    def _build_model(self):
        """Build neural network for Q-learning"""
        model = Sequential([
            Dense(128, activation='relu', input_dim=self.state_size),
            Dropout(0.2),
            Dense(128, activation='relu'),
            Dropout(0.2),
            Dense(64, activation='relu'),
            Dense(self.action_size, activation='linear')
        ])

        model.compile(
            loss='mse',
            optimizer=Adam(learning_rate=self.learning_rate)
        )

        return model

    def update_target_model(self):
        """Copy weights from model to target_model"""
        self.target_model.set_weights(self.model.get_weights())

    def remember(self, state, action, reward, next_state, done):
        """Store experience in replay memory"""
        self.memory.append((state, action, reward, next_state, done))

    def act(self, state):
        """
        Choose action using epsilon-greedy policy
        Returns: action index (0, 1, or 2)
        """
        if np.random.random() <= self.epsilon:
            return random.randrange(self.action_size) # Explore

        q_values = self.model.predict(state, verbose=0)
        return np.argmax(q_values[0]) # Exploit

    def replay(self):
        """Train on batch from replay memory"""
        if len(self.memory) < self.train_start:
            return

```

```

# Sample mini-batch
batch = random.sample(self.memory, min(len(self.memory), self.batch_size))

states = np.array([item[0][0] for item in batch])
actions = np.array([item[1] for item in batch])
rewards = np.array([item[2] for item in batch])
next_states = np.array([item[3][0] for item in batch])
dones = np.array([item[4] for item in batch])

# Predict Q-values
targets = self.model.predict(states, verbose=0)
next_q_values = self.target_model.predict(next_states, verbose=0)

# Update Q-values using Bellman equation
for i in range(len(batch)):
    if dones[i]:
        targets[i][actions[i]] = rewards[i]
    else:
        targets[i][actions[i]] = rewards[i] + self.gamma * np.max(next_q_values[i])

# Train model
self.model.fit(states, targets, epochs=1, verbose=0, batch_size=self.batch_size)

# Decay epsilon
if self.epsilon > self.epsilon_min:
    self.epsilon *= self.epsilon_decay

self.training_step += 1

# Update target model periodically
if self.training_step % 10 == 0:
    self.update_target_model()

def save(self, filepath):
    """Save model to disk"""
    self.model.save(filepath)
    print(f"Model saved to {filepath}")

def load(self, filepath):
    """Load model from disk"""
    if os.path.exists(filepath):
        self.model = load_model(filepath)
        self.update_target_model()
        print(f"Model loaded from {filepath}")
        return True
    return False

def get_metrics(self):
    """Return current agent metrics"""
    return {
        'epsilon': self.epsilon,
        'memory_size': len(self.memory),
        'training_steps': self.training_step
    }

```

Step 2: Create AI Signal Pipeline

Create backend/ai_pipeline.py:

```
import numpy as np
import pandas as pd
from rl_agent import DQNAgent
from datetime import datetime

class AISignalPipeline:
    """
    AI-assisted trading signal generation pipeline
    Combines SMC analysis with RL agent decisions
    """

    def __init__(self, state_size=50, model_path='models/rl_agent.h5'):
        self.state_size = state_size
        self.action_size = 3 # HOLD, BUY, SELL
        self.model_path = model_path

        # Initialize RL agent
        self.agent = DQNAgent(state_size, self.action_size)
        self.agent.load(model_path)

        # Performance tracking
        self.episode_rewards = []
        self.current_episode_reward = 0

    def extract_state(self, df, signal_data):
        """
        Extract state vector from market data and SMC signals
        Returns: numpy array of shape (1, state_size)
        """
        features = []

        # Price-based features
        if len(df) >= 20:
            close_prices = df['close'].values[-20:]

            # Returns
            returns = np.diff(close_prices) / close_prices[:-1]
            features.extend([
                returns.mean(),
                returns.std(),
                returns[-1] if len(returns) > 0 else 0
            ])

            # Price momentum
            features.append((close_prices[-1] - close_prices[0]) / close_prices[0])

            # Volatility
            features.append(close_prices.std() / close_prices.mean())
        else:
            features.extend([0, 0, 0, 0, 0])

        # SMC signal features
```

```

order_blocks = signal_data.get('order_blocks', [])
choch = signal_data.get('choch')
engulfing = signal_data.get('engulfing')

# Order block proximity
if order_blocks:
    latest_ob = order_blocks[-1]
    current_price = df['close'].iloc[-1]

    if latest_ob['type'] == 'BULLISH':
        ob_distance = (current_price - latest_ob['low']) / current_price
        features.append(ob_distance)
        features.append(1.0) # Bullish OB flag
    else:
        ob_distance = (latest_ob['high'] - current_price) / current_price
        features.append(ob_distance)
        features.append(0.0) # Bearish OB flag
else:
    features.extend([0, 0.5])

# CHoCH signal
if choch:
    features.append(1.0 if 'BULLISH' in choch['type'] else -1.0)
else:
    features.append(0.0)

# Engulfing pattern
if engulfing:
    features.append(engulfing['strength'])
    features.append(1.0 if 'BULLISH' in engulfing['type'] else -1.0)
else:
    features.extend([0, 0])

# SMC confidence from original signal
features.append(signal_data.get('confidence', 0.5))

# Pad or truncate to state_size
features = np.array(features)
if len(features) < self.state_size:
    features = np.pad(features, (0, self.state_size - len(features)))
else:
    features = features[:self.state_size]

return features.reshape(1, self.state_size)

def get_action(self, state):
    """Get action from RL agent"""
    action_idx = self.agent.act(state)
    actions = ['HOLD', 'BUY', 'SELL']
    return actions[action_idx], action_idx

def calculate_reward(self, trade_result):
    """
    Calculate reward for RL agent based on trade outcome
    Positive reward for profit, negative for loss
    """

```

```

    if not trade_result:
        return 0

    pnl = trade_result.get('pnl', 0)
    risk = trade_result.get('risk', 1)

    # Normalize reward by risk
    reward = pnl / risk

    # Bonus for maintaining good risk/reward
    if pnl > 0:
        reward += 0.5 # Bonus for winning trade
    elif pnl < 0:
        reward -= 0.3 # Penalty for losing trade

    return reward

def learn_from_trade(self, state, action_idx, trade_result, next_state, done):
    """Update RL agent based on trade outcome"""
    reward = self.calculate_reward(trade_result)

    self.agent.remember(state, action_idx, reward, next_state, done)
    self.agent.replay()

    self.current_episode_reward += reward

    if done:
        self.episode_rewards.append(self.current_episode_reward)
        self.current_episode_reward = 0

def get_signal(self, df, smc_signal):
    """
    Generate enhanced trading signal using AI
    Returns: dict with action, confidence, and state
    """
    state = self.extract_state(df, smc_signal)
    action, action_idx = self.get_action(state)

    # Calculate combined confidence
    smc_confidence = smc_signal.get('confidence', 0.5)

    # Get Q-values for confidence estimation
    q_values = self.agent.model.predict(state, verbose=0)[^0]
    max_q = np.max(q_values)
    min_q = np.min(q_values)
    q_range = max_q - min_q if max_q != min_q else 1

    # Normalize Q-value to 0-1
    ai_confidence = (q_values[action_idx] - min_q) / q_range

    # Combined confidence (weighted average)
    combined_confidence = 0.6 * smc_confidence + 0.4 * ai_confidence

    return {
        'action': action,
        'action_idx': action_idx,
    }

```

```

        'confidence': combined_confidence,
        'ai_confidence': ai_confidence,
        'smc_confidence': smc_confidence,
        'state': state,
        'q_values': q_values.tolist()
    }

def save_model(self):
    """Save RL agent model"""
    self.agent.save(self.model_path)

def get_performance_metrics(self):
    """Get AI pipeline performance metrics"""
    metrics = self.agent.get_metrics()

    if self.episode_rewards:
        metrics['avg_episode_reward'] = np.mean(self.episode_rewards[-100:])
        metrics['total_episodes'] = len(self.episode_rewards)

    return metrics

```

Step 3: Create Enhanced Strategy

Create backend/enhanced_strategy.py:

```

from smc_strategy import SMCStrategy
from ai_pipeline import AISignalPipeline
import pandas as pd

class EnhancedSMCStrategy(SMCStrategy):
    """
    Enhanced SMC Strategy with AI/RL integration
    Combines traditional SMC with reinforcement learning
    """

    def __init__(self, risk_params):
        super().__init__(risk_params)
        self.ai_pipeline = AISignalPipeline()
        self.trade_states = {} # Store states for learning

    def generate_signal(self, df, symbol):
        """
        Generate enhanced signal using both SMC and AI
        """
        # Get base SMC signal
        smc_signal = super().generate_signal(df, symbol)

        if not smc_signal:
            return None

        # Get AI-enhanced signal
        ai_signal = self.ai_pipeline.get_signal(df, smc_signal)

        # Combine signals
        enhanced_signal = {

```

```

        **smc_signal,
        'ai_action': ai_signal['action'],
        'ai_confidence': ai_signal['ai_confidence'],
        'combined_confidence': ai_signal['confidence'],
        'q_values': ai_signal['q_values']
    }

    # Store state for later learning
    self.trade_states[symbol] = {
        'state': ai_signal['state'],
        'action_idx': ai_signal['action_idx']
    }

    # Override direction based on AI if confidence high enough
    if ai_signal['confidence'] > 0.7:
        if ai_signal['action'] == 'BUY':
            enhanced_signal['direction'] = 'BUY'
        elif ai_signal['action'] == 'SELL':
            enhanced_signal['direction'] = 'SELL'
        else: # HOLD
            enhanced_signal['direction'] = 'NEUTRAL'

    return enhanced_signal

def record_trade_outcome(self, symbol, trade_result, next_df):
    """
    Learn from trade outcome
    Called when trade is closed
    """
    if symbol not in self.trade_states:
        return

    prev_state_data = self.trade_states[symbol]

    # Extract next state
    if next_df is not None and len(next_df) > 0:
        smc_signal = super().generate_signal(next_df, symbol)
        if smc_signal:
            next_state = self.ai_pipeline.extract_state(next_df, smc_signal)
        else:
            next_state = prev_state_data['state']
    else:
        next_state = prev_state_data['state']

    # Learn from experience
    self.ai_pipeline.learn_from_trade(
        state=prev_state_data['state'],
        action_idx=prev_state_data['action_idx'],
        trade_result=trade_result,
        next_state=next_state,
        done=True
    )

    # Clean up
    del self.trade_states[symbol]

```

```

def save_models(self):
    """Save AI models"""
    self.ai_pipeline.save_model()

def get_ai_metrics(self):
    """Get AI performance metrics"""
    return self.ai_pipeline.get_performance_metrics()

```

Part 2: Continuous Learning Pipeline

Create Training Pipeline

Create backend/training_pipeline.py:

```

import schedule
import time
from datetime import datetime, timedelta
import pandas as pd
import numpy as np

class ContinuousLearningPipeline:
    """
    Continuous model retraining pipeline
    Automatically improves models based on recent trading data
    """

    def __init__(self, strategy, delta_client):
        self.strategy = strategy
        self.delta_client = delta_client
        self.training_interval_hours = 24
        self.last_training = datetime.now()
        self.min_trades_for_training = 20

    def collect_recent_trades(self, hours=24):
        """Collect recent closed trades for training"""
        trades = self.strategy.closed_trades

        if not trades:
            return []

        # Filter trades from last N hours
        cutoff_time = datetime.now() - timedelta(hours=hours)

        recent_trades = [
            t for t in trades
            if datetime.fromisoformat(t.get('closed_at', '')) > cutoff_time
        ]

        return recent_trades

    def prepare_training_data(self, trades):
        """
        Prepare training data from trade history
        """

```

```

    Returns: features (X) and labels (y)
    """
    X = []
    y = []

    for trade in trades:
        if 'features' not in trade:
            continue

        features = trade['features']
        # Label: 1 for profitable, 0 for loss
        label = 1 if trade.get('pnl', 0) > 0 else 0

        X.append(features)
        y.append(label)

    return np.array(X), np.array(y)

def retrain_models(self):
    """
    Retrain RL agent based on recent performance
    """
    print(f"\n{'='*50}")
    print("Starting continuous learning session")
    print(f"{'='*50}\n")

    # Collect recent trades
    recent_trades = self.collect_recent_trades(hours=self.training_interval_hours)

    print(f" Collected {len(recent_trades)} recent trades")

    if len(recent_trades) < self.min_trades_for_training:
        print(f"⚠ Insufficient trades ({len(recent_trades)} < {self.min_trades_f
        print(" Skipping retraining")
        return

    # Calculate performance metrics
    profitable_trades = sum(1 for t in recent_trades if t.get('pnl', 0) > 0)
    win_rate = profitable_trades / len(recent_trades)
    avg_pnl = np.mean([t.get('pnl', 0) for t in recent_trades])

    print(f" Performance Summary:")
    print(f" Win Rate: {win_rate*100:.1f}%")
    print(f" Avg PnL: ${avg_pnl:.2f}")

    # Retrain RL agent
    print("\n Retraining RL agent...")

    for trade in recent_trades:
        if 'state' in trade and 'action_idx' in trade:
            reward = trade.get('pnl', 0) / max(trade.get('risk', 1), 0.01)

            self.strategy.ai_pipeline.agent.remember(
                state=trade['state'],
                action=trade['action_idx'],
                reward=reward,

```

```

        next_state=trade.get('next_state', trade['state']),
        done=True
    )

    # Perform multiple replay sessions
    for _ in range(10):
        self.strategy.ai_pipeline.agent.replay()

    # Save updated models
    self.strategy.save_models()

    self.last_training = datetime.now()

    print(f"\n✓ Retraining complete at {self.last_training}")
    print(f"  Next training in {self.training_interval_hours} hours")
    print(f"{'='*50}\n")

def schedule_training(self):
    """Schedule periodic retraining"""
    schedule.every(self.training_interval_hours).hours.do(self.retrain_models)

    print(f"⌚ Scheduled continuous learning every {self.training_interval_hours} hour")

    while True:
        schedule.run_pending()
        time.sleep(3600) # Check every hour

def manual_retrain(self):
    """Trigger manual retraining"""
    print("⌚ Manual retraining triggered")
    self.retrain_models()

```

Part 3: Google Colab Training

Create Colab Notebook

Create training/google_colab_training.ipynb:

```

# SMC Trading Bot - Model Training on Google Colab
# Run this notebook to train models using free GPU

# Mount Google Drive
from google.colab import drive
drive.mount('/content/drive')

# Install dependencies
!pip install tensorflow pandas numpy scikit-learn

# Download historical data
import requests
import pandas as pd
import numpy as np

```

```

def fetch_training_data(symbol='BTCUSD', days=365):
    """Fetch historical data for training"""
    # Use public API to get historical data
    url = f"https://api.delta.exchange/v2/history/candles"

    end_time = int(time.time())
    start_time = end_time - (days * 24 * 60 * 60)

    params = {
        'symbol': symbol,
        'resolution': '15',
        'start': start_time,
        'end': end_time
    }

    response = requests.get(url, params=params)
    data = response.json()

    df = pd.DataFrame(
        data['result'],
        columns=['time', 'open', 'high', 'low', 'close', 'volume']
    )

    return df

# Fetch data for BTC and ETH
print("Fetching training data...")
btc_data = fetch_training_data('BTCUSD', days=365)
eth_data = fetch_training_data('ETHUSD', days=365)

print(f"BTC data: {len(btc_data)} candles")
print(f"ETH data: {len(eth_data)} candles")

# Train RL agent
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.optimizers import Adam

def build_model(state_size=50, action_size=3):
    model = Sequential([
        Dense(128, activation='relu', input_dim=state_size),
        Dropout(0.2),
        Dense(128, activation='relu'),
        Dropout(0.2),
        Dense(64, activation='relu'),
        Dense(action_size, activation='linear')
    ])

    model.compile(
        loss='mse',
        optimizer=Adam(learning_rate=0.001)
    )

    return model

model = build_model()

```

```

# Generate training samples
# (Add your feature extraction logic here)

# Train model
print("Training model...")
history = model.fit(
    X_train, y_train,
    epochs=50,
    batch_size=32,
    validation_split=0.2,
    verbose=1
)

# Save model
model.save('/content/drive/MyDrive/smc_models/rl_agent.h5')
print("Model saved to Google Drive!")

# Plot training history
import matplotlib.pyplot as plt

plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.legend()
plt.title('Model Training History')
plt.show()

```

Part 4: Integration with Main Bot

Update backend/app.py

Add RL integration:

```

# Import enhanced strategy
from enhanced_strategy import EnhancedSMCStrategy
from training_pipeline import ContinuousLearningPipeline

# Replace SMC strategy with enhanced version
smc_strategy = EnhancedSMCStrategy(RISK_PARAMS)

# Initialize training pipeline
training_pipeline = ContinuousLearningPipeline(smc_strategy, delta_client)

# Add new API endpoints
@app.route('/api/ai/metrics', methods=['GET'])
def get_ai_metrics():
    """Get AI/RL performance metrics"""
    metrics = smc_strategy.get_ai_metrics()
    return jsonify(metrics)

@app.route('/api/ai/retrain', methods=['POST'])
def trigger_retrain():
    """Manually trigger model retraining"""

```

```

    training_pipeline.manual_retrain()
    return jsonify({'status': 'Retraining initiated'})

# Start training pipeline in background
import threading
training_thread = threading.Thread(
    target=training_pipeline.schedule_training,
    daemon=True
)
training_thread.start()

```

Part 5: Enhanced Dashboard

Update Dashboard with AI Metrics

Add to frontend/static/js/app.js:

```

// Fetch AI metrics
async function fetchAIMetrics() {
    const response = await fetch(`.${API_BASE}/ai/metrics`);
    return await response.json();
}

// Update AI metrics display
async function updateAIMetrics() {
    const metrics = await fetchAIMetrics();

    document.getElementById('ai-epsilon').textContent =
        metrics.epsilon.toFixed(3);
    document.getElementById('ai-memory-size').textContent =
        metrics.memory_size;
    document.getElementById('ai-training-steps').textContent =
        metrics.training_steps;

    if (metrics.avg_episode_reward) {
        document.getElementById('ai-avg-reward').textContent =
            metrics.avg_episode_reward.toFixed(2);
    }
}

// Trigger manual retraining
async function triggerRetraining() {
    const response = await fetch(`.${API_BASE}/ai/retrain`, {
        method: 'POST'
    });
    const result = await response.json();
    alert(result.status);
}

// Update every 10 seconds
setInterval(updateAIMetrics, 10000);

```

Add HTML for AI metrics:

```
<div>
  <h3>AI Performance</h3>

  <div>
    <div>
      <p>Exploration Rate</p>
      <p>0.500</p>
    </div>

    <div>
      <p>Memory Size</p>
      <p>0</p>
    </div>

    <div>
      <p>Training Steps</p>
      <p>0</p>
    </div>

    <div>
      <p>Avg Reward</p>
      <p>0.00</p>
    </div>
  </div>

  &lt;button onclick="triggerRetraining()" 
            class="mt-4 w-full bg-purple-600 hover:bg-purple-700 
            text-white px-4 py-2 rounded"&gt;
    Retrain Now
  &lt;/button&gt;
</div>
```

Part 6: Deployment

Update Requirements

Add to backend/requirements.txt:

```
tensorflow==2.15.0
```

Rebuild and Deploy

```
# Local
docker-compose down
docker-compose build --no-cache
docker-compose up -d
```

```
# VPS (via git push - GitHub Actions will deploy)
git add .
git commit -m "Add Phase 3: RL/AI integration"
git push origin main
```

Part 7: Monitoring AI Performance

Track RL Agent Metrics

Monitor:

- **Epsilon:** Exploration vs exploitation balance
- **Memory Size:** Experiences stored
- **Training Steps:** Number of training iterations
- **Average Reward:** Performance indicator

Expected Behavior

- Epsilon starts at 1.0 (full exploration)
- Gradually decreases to 0.01 (mostly exploitation)
- Memory fills up to 2000 experiences
- Average reward should increase over time

Summary

- ✓ **Phase 3 Complete** - AI/RL integrated
- ✓ DQN agent learning from trades
- ✓ Continuous learning pipeline active
- ✓ Google Colab ready for GPU training
- ✓ Enhanced dashboard with AI metrics

Free Services Used:

- TensorFlow (open-source)
- Google Colab (free GPU)
- All running on Oracle Cloud free tier

Next: Monitor AI performance and fine-tune hyperparameters!

[⁷¹][⁷²][⁷⁵]

[\[1\]](#) [\[2\]](#) [\[3\]](#) [\[4\]](#) [\[5\]](#) [\[6\]](#) [\[7\]](#) [\[8\]](#) [\[9\]](#) [\[10\]](#) [\[11\]](#) [\[12\]](#) [\[13\]](#) [\[14\]](#) [\[15\]](#) [\[16\]](#)

**

1. https://help.ovhcloud.com/csm/en-vps-deploy-website-github-actions?id=kb_article_view&sysparm_article=KB0066184
2. <https://deltaexchangeindia.freshdesk.com/support/solutions/articles/80001181136-getting-started-with-delta-exchange-api-copilot>
3. <https://docs.docker.com/reference/compose-file/services/>
4. <https://www.hostinger.com/support/deploy-to-hostinger-vps-using-github-actions/>
5. <https://www.youtube.com/watch?v=TSzz-Fllq14>
6. <https://learn.microsoft.com/en-us/dotnet/architecture/microservices/multi-container-microservice-net-application/multi-container-applications-docker-compose>
7. <https://blog.ando.ai/posts/github-actions-vps-deployment/>
8. <https://www.youtube.com/watch?v=O5i5DJUAap4>
9. <https://docs.docker.com/compose/how-tos/multiple-compose-files/>
10. <https://docs.github.com/actions/deployment/about-deployments/deploying-with-github-actions>
11. <https://www.delta.exchange/support/solutions/articles/80001153884-your-complete-api-faq-guide-everything-you-need-to-know>
12. <https://www.profitaddaweb.com/2025/04/delta-exchange-api-in-python.html>
13. <https://stackoverflow.com/questions/49191304/how-to-deal-with-multiple-services-inside-a-docker-compose-yml>
14. <https://gist.github.com/danielwetan/4f4db933531db5dd1af2e69ec8d54d8a>
15. <https://www.youtube.com/watch?v=RoVeASkeMpc>
16. <https://matthiasnoback.nl/2018/03/defining-multiple-similar-services-with-docker-compose/>